

ECB Project

Contents

| | |
|--|-----------|
| Preliminaries | 2 |
| Setup | 2 |
| Interactive Option Selection | 2 |
| | |
| Data | 4 |
| Sources & Explanations | 4 |
| Loading & Preparation Data | 4 |
| Options Configuration | 6 |
| Raw Data Plots | 6 |
| Data Properties | 7 |
| | |
| Taylor Rule Estimation | 10 |
| Without Lags | 10 |
| Lagged Models | 10 |
| Checking for structural breaks | 11 |
| Rolling Estimation (for structural breaks) | 14 |
| | |
| Forecasting Model Evaluation | 19 |
| Helpers | 19 |
| Estimation | 22 |
| Spaghetti Plots | 26 |
| Plots of FE | 27 |
| Density of FE | 28 |
| Variance of FE | 31 |
| Absolute Performance: Efficiency & Bias | 32 |
| Relative Performance (against benchmark) | 35 |
| | |
| Actual Forecast Model | 40 |
| Helpers | 40 |
| Forecasting | 42 |
| | |
| Prediction Intervals | 44 |

Preliminaries

Setup

```
# Clear memory
rm(list=ls())

# Load here package to enable finding script loading other packages
require(here)

# Set directory
getwd()
setwd("...")

# Load packages & helper functions
source(here("helpers/packages.R"))
source(here("helpers/auto_ARIMA_replic.R"))

# Api key for data
fredr_set_key("e0169694a62c1337f1969e3872605eca")

# Dates (to automatically get the latest data from API calls)
start_date <- "1999-01-01"
end_date <- Sys.Date()

# For replication
set.seed(2025)
```

Interactive Option Selection

- Use Hamilton Filter:
 - TRUE: Selects Hamilton method for output gap estimation
 - FALSE: Selects Hodrick-Prescott method for output gap estimation
- Use Inflation Expectations:
 - TRUE: The models used for forecasting will use 12-month ahead inflation expectations from the ECB survey of professional forecasts (average).
 - FALSE: The models used for forecasting will use realised inflation
- Use Formula
 - Formula 1: Actual interest rate regressed on inflation and output gaps
 - Formula 2: Shadow interest rate regressed on inflation and output gaps
 - Formula 3: Actual interest rate regressed on the one-quarter lag of the interest rate and on inflation and output gaps
 - Formula 4: Shadow interest rate regressed on the one-quarter lag of the shadow interest rate and on inflation and output gaps
- Format:
 - html: For outputting in console or knitting to html
 - latex: For knitting to pdf

```
# Related to Analysis
USE_HAMILTON_FILTER <- TRUE
USE_INFLATION_EXPECTATIONS <- FALSE
USE_FORMULA <- "Formula 3"

# Related to Document Output
format <- "html"
format <- "latex"
```

Data

Sources & Explanations

- FRED Data Source: European Central Bank, ECB Deposit Facility Rate for Euro Area [ECBDFR], retrieved from FRED, Federal Reserve Bank of St. Louis; <https://fred.stlouisfed.org/series/ECBDFR>. The data is the Deposit Facility and is directly reprinted from the ECB. It's in Percent and not seasonally adjusted. The ECB Monetary Policy is steered through this rate
- Shadow Interest Rate: The Shadow rate was developed by **wuTimeVaryingLowerBound2017** and does quantify a hypothetical removal of the ZLB. The rate does not track 1:1 on the deposit facility in the data before the ZLB, instead being closer to the refinancing rate. (maybe need to adjust)
- GDP: The GDP Data is quarterly real GDP in 2010 Euros in million retrieved from FRED via Eurostat. The data is seasonally adjusted.
- Inflation:
 - Realised: The Inflation data is from Eurostat and called HICP. It's the index the ECB uses for Inflation and is not seasonally adjusted.
 - Expectations: Expected Inflation is the ECB's survey of professional forecasters. It forecasts the HICP 12 months in advance.

Loading & Preparation Data

```
# --- 1. ECB Deposit Facility Rate & Shadow Rate ---
ecb_rate_daily <- fredr(series_id = "ECBDFR", observation_start = as.Date(start_date))
ecb_rate_q <- ecb_rate_daily %>%
  mutate(quarter = as.yearqtr(date)) %>%
  group_by(quarter) %>%
  summarise(rate = last(value)) %>%
  mutate(date = as.Date(quarter))
# Wu-Xia Shadow Rate
shadow_rate_daily = as.data.frame(readMat("data/shadowrate_ECB.mat"))
colnames(shadow_rate_daily) <- c("DATE", "shadowrate")
shadow_rate_daily$DATE <- as.Date(paste0(shadow_rate_daily$DATE, "01"), format="%Y%m%d")
shadow_rate_daily$quarter <- as.yearqtr(as.Date(shadow_rate_daily$DATE))
shadow_rate_daily$month <- as.yearmon(as.Date(shadow_rate_daily$DATE))
quarterly_shadow = aggregate(shadowrate ~ quarter, data=shadow_rate_daily, FUN=mean, na.rm=T)
monthly_shadow = aggregate(shadowrate ~ month, data=shadow_rate_daily, FUN=mean, na.rm=T)

# --- 2. HICP Inflation (Euro Area) ---
inflation_data <- get_eurostat("prc_hicp_manr", filters = list(geo = "EA", coicop = "CP00"), type = "la
inflation_q <- inflation_data %>%
  filter(time >= start_date) %>%
  dplyr::select(date = time, inflation = values) %>%
  mutate(quarter = as.yearqtr(date)) %>%
  group_by(quarter) %>%
  summarise(inflation = mean(inflation, na.rm = TRUE)) %>%
  mutate(date = as.Date(quarter))

#inflation expectations
inflation_exp <- rdb(ids = "ECB/SPF/M.U2.HICP.POINT.P12M.Q.AVG")
```

```

#inflation_exp <- rdb(ids = "ECB/SPF/M.U2.HICP.POINT.P24M.Q.AVG")
inflation_exp_q <- inflation_exp %>%
  mutate(quarter = as.yearqtr(period)) %>%
  group_by(quarter) %>%
  summarise(exp_inflation = last(original_value)) %>%
  mutate(date = as.Date(quarter))

#P12M : 12-month ahead forecasts
inflation_q$exp_inflation = c(rep(NA,3),as.numeric(inflation_exp_q$exp_inflation),NA)
#P24M : 24-month ahead forecasts
#inflation_q$exp_inflation = c(rep(NA,7),as.numeric(inflation_exp_q$exp_inflation[1:101]))

# --- 3. Real GDP and Estimated Output Gap ---
# a) Real GDP for the Euro Area. The series ID is CLVMNACSCAB1GQE_A.
gdp_q <- fredr(
  series_id = "CLVMEURSCAB1GQEA19",
  observation_start = as.Date(start_date)) %>%
  mutate(quarter = as.yearqtr(date)) %>%
  dplyr::select(quarter, real_gdp = value) %>%
  mutate(log_real_gdp = log(real_gdp))

# b) Estimate Potential GDP (the trend) using the HP Filter on the log of real GDP.
# The lambda value of 1600 is standard for quarterly data.
hp_gdp <- hpfilter(gdp_q$log_real_gdp, freq = 1600)
gdp_q$potential_gdp_log <- as.numeric(hp_gdp$trend)
ham_gdp_cycle <- filter_hamilton(gdp_q$log_real_gdp, p = 4, horizon = 8)
gdp_q$potential_gdp_log_ham <- gdp_q$log_real_gdp - ham_gdp_cycle

# Combine all data into a single data frame
data <- ecb_rate_q %>%
  dplyr::select(quarter, rate) %>%
  left_join(inflation_q, by = "quarter") %>%
  left_join(gdp_q, by = "quarter") %>%
  left_join(quarterly_shadow, by = "quarter")

# Create model variables
data <- data %>%
  mutate(
    realised_inflation_gap = inflation - 2.0,
    exp_inflation_gap = exp_inflation -2.0,
    output_gap_hp = 100 * (log_real_gdp - potential_gdp_log),
    output_gap_ham = 100 * (log_real_gdp - potential_gdp_log_ham),
    rate_lag = lag(rate, 1),
    shadowrate = case_when(
      quarter < "2012 Q3" | quarter >= "2022 Q3" ~ rate,
      TRUE ~ shadowrate),
    shadowrate_lag = lag(shadowrate, 1))

# Remove last row since no output data
data = subset(data, quarter < "2025 Q3")

# Clean environment
rm(gdp_q, hp_gdp, ecb_rate_daily, ecb_rate_q, inflation_data, inflation_q,

```

```
inflation_exp, inflation_exp_q, monthly_shadow, quarterly_shadow, shadow_rate_daily)
```

Options Configuration

```
# Choices in setup chunk

# ----- 1. Filter selection for output gap estimation -----

# TRUE = Use Hamilton Filter (newer, arguably more robust)
# FALSE = Use HP Filter (classic approach)

# Applying selection
if (USE_HAMILTON_FILTER) {
  data$output_gap <- data$output_gap_ham
  cat(">> CONFIGURATION: Using Hamilton Filter for output gap estimation.")
} else {
  data$output_gap <- data$output_gap_hp
  cat(">> CONFIGURATION: Using HP Filter for output gap estimation.") }
```

CONFIGURATION: Using Hamilton Filter for output gap estimation.

```
# ----- 2. Inflation expectations choice -----

# TRUE = Use inflation expectations from ECB survey of professional forecasts
# FALSE = Use realised inflation

# Applying selection
if (USE_INFLATION_EXPECTATIONS) {
  data$inflation_gap <- data$exp_inflation_gap
  cat("* CONFIGURATION: Using inflation expectations in Taylor Rule forecasting.")
} else {
  data$inflation_gap <- data$realised_inflation_gap
  cat("* CONFIGURATION: Using realised inflation in Taylor Rule forecasting.") }
```

- CONFIGURATION: Using realised inflation in Taylor Rule forecasting.

Raw Data Plots

```
# Data must be in long format for a faceted plot
plot_data <- data %>%
  pivot_longer(cols = c(rate, inflation, exp_inflation, output_gap_hp, output_gap_ham),
              names_to = "series",
              values_to = "value") %>%
  mutate(series = factor(series,
                        levels = c("rate", "inflation", "exp_inflation", "output_gap_hp", "output_gap_ham"),
                        labels = c("Deposit Rate", "(Realised) Inflation", "(Expected) Inflation", "Output Gap (HP)", "Output Gap (Hamilton)")))

ggplot(plot_data, aes(x = date, y = value)) +
```


```

geom_line(aes(color = series), linewidth = 1) +
facet_wrap(~ series, scales = "free_y", ncol = 2) +
scale_x_date(date_breaks = "5 years", date_labels = "%Y") +
labs(title = "Raw Data Plots", subtitle = "Y axis in %", x = "", y = "") +
theme_minimal() + theme(
  plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
  plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
  axis.title = element_text(size = 12),
  axis.text = element_text(size = 10),
  legend.position = "none",
  panel.grid.major.x = element_line(linewidth = 0.525, color = "grey83"))

```

Raw Data Plots

Y axis in %



```
rm(plot_data)
```

Data Properties

```

# Helper function to run ADF and KPSS tests and return dataframe fit for table
check_stationarity <- function(y, var_name) {
  # Running tests and getting pvalues
  adf_result <- tseries::adf.test(y, alternative = "stationary")
  kpss_result <- tseries::kpss.test(y, null = "Level")

```

```

adf_p <- adf_result$p.value
kpss_p <- kpss_result$p.value

# Automatically assign interpretation (yes, modular!)
interpretation <- ""
if (adf_p < 0.05 && kpss_p > 0.05) {
  interpretation <- "Stationary I(0): ADF rejects unit root, KPSS confirms stationarity. "
} else if (adf_p > 0.05 && kpss_p < 0.05) {
  interpretation <- "Non-Stationary I(1): ADF confirms unit root, KPSS rejects stationarity. "
} else if (adf_p > 0.05 && kpss_p > 0.05) {
  interpretation <- "Conflicting Results: ADF confirms unit root, while KPSS suggests stationarity. "
} else { # adf_p < 0.05 && kpss_p < 0.05
  interpretation <- "Conflicting Results: ADF rejects unit root, while KPSS rejects stationarity. "
}

# Export results
data.frame(Variable = var_name,
           `ADF p-value` = adf_p,
           `KPSS p-value` = kpss_p,
           Result = interpretation, check.names = FALSE) }

# Similar helper function, but for cointegration tests
check_coint <- function(y1,y2, var_name1, var_name2){

  # Run test and get pvalue
  coint_test <- aTSA::coint.test(y1, y2, output = FALSE)
  p_value <- coint_test[1, 3] #type 1

  # Again, modular(!) interpretation
  interpretation <- ""
  if (p_value < 0.05) {
    interpretation <- "Cointegrated"
  } else {
    interpretation <- "Not Cointegrated"
  }

  # Export results
  data.frame(Variables = paste(var_name1, "&", var_name2),
             `p-value` = p_value,
             Result = interpretation, check.names = FALSE) }

# ----- Run Tests -----

# Use helper to run stationarity checks for our main variables
test_rate <- check_stationarity(data$rate, "Interest Rate")
test_inflation <- check_stationarity(data$inflation, "Inflation")
test_output_gap <- check_stationarity(na.omit(data$output_gap), "Output Gap")

# Given results for rate and inflation, run tests on 1st diffs
test_rate_diff <- check_stationarity(diff(data$rate), "Interest Rate (1st Diff)")
test_inflation_diff <- check_stationarity(diff(data$inflation), "Inflation (1st Diff)")

# Since rate and inflation are I(1), run a cointegration test
coint_test = check_coint(data$rate, data$inflation,

```

Table 1: Summary of Stationarity Tests (ADF & KPSS)

| Variable | ADF p-value | KPSS p-value | Result |
|--------------------------|-------------|--------------|--|
| Interest Rate | 0.2900 | 0.0325 | Non-Stationary I(1): ADF confirms unit root, KPSS rejects stationarity. |
| Inflation | 0.2319 | 0.1000 | Conflicting Results: ADF confirms unit root, while KPSS suggests stationarity. |
| Output Gap | 0.0144 | 0.1000 | Stationary I(0): ADF rejects unit root, KPSS confirms stationarity. |
| Interest Rate (1st Diff) | 0.0100 | 0.1000 | Stationary I(0): ADF rejects unit root, KPSS confirms stationarity. |
| Inflation (1st Diff) | 0.0100 | 0.1000 | Stationary I(0): ADF rejects unit root, KPSS confirms stationarity. |

Table 2: Cointegration Test Results

| Variables | p-value | Result |
|---------------------------|---------|--------------|
| Interest Rate & Inflation | 0.0433 | Cointegrated |

```

var_name1 = "Interest Rate", var_name2 = "Inflation")

# ----- Print Results -----

# Combine stationarity results
all_stationarity_results <- rbind(test_rate,
                                    test_inflation,
                                    test_output_gap,
                                    test_rate_diff,
                                    test_inflation_diff)

# Tables
kable(all_stationarity_results, digits = 4, format = format, booktabs = TRUE,
      caption = "Summary of Stationarity Tests (ADF \& KPSS)") %>%
  kable_styling(latex_options = "scale_down",
                position = "center") %>%
  column_spec(1, border_right = TRUE) %>%
  column_spec(4, width = "6cm")

kable(coint_test, digits = 4, format = format, booktabs = TRUE,
      caption = "Cointegration Test Results") %>%
  kable_styling(latex_options = "scale_down",
                position = "center") %>%
  column_spec(1, border_right = TRUE)

```

Taylor Rule Estimation

Without Lags

$$i_t = \pi^* + \beta(\pi_t - \pi^*) + \gamma(y_t - \bar{y}_t)$$

```
TR <- lm(rate ~ realised_inflation_gap + output_gap, data = data)
TRsr <- lm(shadowrate ~ realised_inflation_gap + output_gap, data = data)

table1 <- export_summs(TR, TRsr, vcov = sandwich::NeweyWest,
                       model.names = c("TR", "TR w/ SR"), digits = 4)
huxtable::caption(table1) <- "No Lag, No Expectations"
table1
```

Table 3: No Lag, No Expectations

| | TR | TR w/ SR |
|------------------------|----------|----------|
| (Intercept) | 0.8345 | -0.8601 |
| | (0.9055) | (2.6782) |
| realised_inflation_gap | 0.1963 | 0.6857 |
| | (0.4289) | (0.6062) |
| output_gap | 0.0824 | -0.0567 |
| | (0.1628) | (0.2758) |
| N | 95 | 95 |
| R2 | 0.1726 | 0.1159 |

*** p < 0.001; ** p < 0.01; * p < 0.05.

```
TR_e <- lm(rate ~ exp_inflation_gap + output_gap, data = data)
TRsr_e <- lm(shadowrate ~ exp_inflation_gap + output_gap, data = data)
TR_ie <- lm(rate ~ realised_inflation_gap + exp_inflation_gap + output_gap, data = data)
TRsr_ie <- lm(shadowrate ~ realised_inflation_gap + exp_inflation_gap + output_gap, data = data)

table2 <- export_summs(TR_e, TRsr_e, TR_ie, TRsr_ie, vcov = sandwich::NeweyWest,
                       model.names = c("TR", "TR w/ SR", "TR", "TR w/ SR"), digits = 4)
huxtable::caption(table2) <- "No Lag, with Inflation Expectations"
table2
```

Lagged Models

$$i_t = \pi^* + \phi i_{t-1} + \beta(\pi_t - \pi^*) + \gamma(y_t - \bar{y}_t)$$

Table 4: No Lag, with Inflation Expectations

| | TR | TR w/ SR | TR | TR w/ SR |
|------------------------|------------------------|-----------------------|------------------------|-----------------------|
| (Intercept) | 1.3440 *** (0.3179) | 0.2867 (1.5603) | 1.3381 *** (0.3714) | 0.1715 (1.5918) |
| exp_inflation_gap | 1.7238 *** (0.3523) | 3.7581 ** (1.3756) | 1.7107 *** (0.3466) | 3.5044 ** (1.1674) |
| output_gap | 0.0707 (0.0468) | -0.0026 (0.1871) | 0.0662 (0.0810) | -0.0898 (0.2133) |
| realised_inflation_gap | | | 0.0163 (0.1277) | 0.3171 (0.3433) |
| N | 95 | 95 | 95 | 95 |
| R2 | 0.6177 | 0.3904 | 0.6180 | 0.4086 |

*** p < 0.001; ** p < 0.01; * p < 0.05.

```

lTR <- lm(rate ~ rate_lag + realised_inflation_gap + output_gap, data = data)
lTRsr <- lm(shadowrate ~ shadowrate_lag + realised_inflation_gap + output_gap, data = data)

table3 <- export_summs(lTR, lTRsr, vcov = sandwich::NeweyWest,
  model.names = c("TR", "TR w/ SR"), digits = 4)
huxtable::caption(table3) <- "Interest Rate Lag, No Expectations"
table3

lTR_e <- lm(rate ~ rate_lag + exp_inflation_gap + output_gap, data = data)
lTRsr_e <- lm(shadowrate ~ shadowrate_lag + exp_inflation_gap + output_gap, data = data)
lTR_ie <- lm(rate ~ rate_lag + realised_inflation_gap + exp_inflation_gap + output_gap, data = data)
lTRsr_ie <- lm(shadowrate ~ shadowrate_lag + realised_inflation_gap + exp_inflation_gap + output_gap, d

table4 <- export_summs(lTR_e, lTRsr_e, lTR_ie, lTRsr_ie, vcov = sandwich::NeweyWest,
  model.names = c("TR", "TR w/ SR", "TR", "TR w/ SR"), digits = 4)
huxtable::caption(table4) <- "Interest Rate Lag, with Inflation Expectations"
table4

```

Checking for structural breaks

```

# Formula to test for breaks
break_formula = rate ~ rate_lag + inflation_gap + output_gap

# Suspected break: Start of ZLB in 2012 Q3
breakpoint1_obs = 55 #R = 55 in evaluation chunk

```

Table 5: Interest Rate Lag, No Expectations

| | TR | TR w/ SR |
|------------------------|------------------------|------------------------|
| (Intercept) | 0.0507 (0.0410) | -0.0777 (0.0618) |
| rate_lag | 0.9189 *** (0.0390) | |
| realised_inflation_gap | 0.0879 * (0.0350) | 0.2498 *** (0.0344) |
| output_gap | 0.0218 (0.0163) | -0.0116 (0.0185) |
| shadowrate_lag | | 0.9530 *** (0.0176) |
| N | 95 | 95 |
| R2 | 0.9594 | 0.9828 |

*** p < 0.001; ** p < 0.01; * p < 0.05.

```

breakpoint1_date <- data$quarter[breakpoint1_obs]

# Suspected break: Covid
breakpoint2_obs = 85 #R = 85 in evaluation chunk
breakpoint2_date <- data$quarter[breakpoint2_obs]

# Chow test (rejecting the null means there are structural breaks)
chow_test1 <- sctest(break_formula, type = "Chow", point = breakpoint1_obs, data = data)
chow_test2 <- sctest(break_formula, type = "Chow", point = breakpoint2_obs, data = data)

# Table with Chow results (for suspected breaks)
chow_df <- data.frame(
  Event = c("ZLB Start", "COVID-19 Start"),
  Date = c(as.character(breakpoint1_date), as.character(breakpoint2_date)),
  `p-value` = c(chow_test1$p.value, chow_test2$p.value), check.names = FALSE)

kable(chow_df, digits = 4, format = format, booktabs = TRUE,
      caption = "Chow tests for suspected structural breaks") %>%
  kable_styling(latex_options = "scale_down",
               position = "center") %>%
  column_spec(1, border_right = TRUE)

```

Table 6: Interest Rate Lag, with Inflation Expectations

| | TR | TR w/ SR | TR | TR w/ SR |
|------------------------|------------|------------|------------|------------|
| (Intercept) | 0.1804 * | 0.0009 | 0.1341 *** | -0.0896 |
| | (0.0796) | (0.0869) | (0.0378) | (0.0622) |
| rate_lag | 0.8606 *** | | 0.8742 *** | |
| | (0.0438) | | (0.0392) | |
| exp_inflation_gap | 0.2387 ** | 0.1296 | 0.1536 *** | -0.0546 |
| | (0.0733) | (0.1039) | (0.0435) | (0.0622) |
| output_gap | 0.0449 * | 0.0592 | 0.0233 | -0.0108 |
| | (0.0218) | (0.0474) | (0.0165) | (0.0182) |
| shadowrate_lag | | 0.9630 *** | | 0.9581 *** |
| | | (0.0287) | | (0.0190) |
| realised_inflation_gap | | | 0.0770 * | 0.2532 *** |
| | | | (0.0378) | (0.0347) |
| N | 95 | 95 | 95 | 95 |
| R2 | 0.9544 | 0.9712 | 0.9611 | 0.9828 |

*** p < 0.001; ** p < 0.01; * p < 0.05.

```

# Estimate Bai-Perron test & output results
BP_test = breakpoints(break_formula, data = data)
BP_test_res = summary(BP_test)

# Optimal values (modular!)
bic_values <- BP_test_res$RSS[2, ]
optimal_m <- as.numeric(names(bic_values)[which.min(bic_values)])

# Make a table out of results (also modular!)
if (optimal_m == 0) {
  bp_df <- data.frame(
    `Detected Breaks` = "No structural breaks detected",
    check.names = FALSE)
} else {
  
```

Table 7: Chow tests for suspected structural breaks

| Event | Date | p-value |
|----------------|---------|---------|
| ZLB Start | 2012 Q3 | 0.0020 |
| COVID-19 Start | 2020 Q1 | 0.0365 |

Table 8: Bai-Perron test for multiple breaks

| Detected Breaks |
|-----------------|
| 2006 Q2 |
| 2019 Q1 |

```

break_obs <- na.omit(BP_test_res$breakpoints[optimal_m, ])
detected_dates <- data$quarter[break_obs]
bp_df <- data.frame(
  `Detected Breaks` = as.character(detected_dates),
  check.names = FALSE) }

kable(bp_df, format = format, booktabs=TRUE,
  caption = "Bai-Perron test for multiple breaks") %>%
  kable_styling(latex_options = "scale_down",
    position = "center")

```

Note: *breaks_obs* shows in which row the BP breaks are

Rolling Estimation (for structural breaks)

```

# This function automatically plots the output from the rolling estimation loop
plot_rolling_coefs <- function(data, var_name, var_name_title = var_name) {

  # 1. Dynamically create the column names for CIs
  lower_col <- paste0(var_name, "_lower")
  upper_col <- paste0(var_name, "_upper")

  # 2. Create the plot, using .data[][] to find the
  #      columns based on the variable names (modularity)
  plot <- ggplot(data, aes(x = date)) +

    # Add a dashed line at y=0 for reference
    geom_hline(yintercept = 0, linetype = "dashed", color = "grey40", linewidth = 0.5) +

    # Add 95% confidence interval ribbon
    geom_ribbon(aes(ymin = .data[[lower_col]], ymax = .data[[upper_col]]),
      fill = "dodgerblue", alpha = 0.3) +

    # Add the coefficient estimate line
    geom_line(aes(y = .data[[var_name]]),
      color = "dodgerblue4", linewidth = 1) +

    # Aesthetic
    labs(title = paste("Rolling Coefficient Estimate:", var_name_title),
      subtitle = "with 95% confidence interval",
      x = "",
      y = "Coefficient Value") +
    theme_minimal() +
}

```

```

theme(plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
      plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
      axis.title = element_text(size = 12),
      axis.text = element_text(size = 10))
return(plot)
}

# Set rolling window (in quarters) & Looping Parameter
W = 30
L = nrow(data) - W + 1
formula = rate ~ rate_lag + inflation_gap + output_gap

# Preparation of result data, dates, var names, and confidence intervals
var_names <- attr(terms(formula), "term.labels")
window_end_dates <- data$quarter[W:nrow(data)] # First window [1:W] ends at data$date[W]
TR_roll <- data.frame(date = window_end_dates)
TR_roll[var_names] <- NA
lower_col_names <- paste0(var_names, "_lower")
upper_col_names <- paste0(var_names, "_upper")

# Looped estimation of TR, outputs coefficients and CIs
for (l in 1:L) {
  # 1. Define splits (with rolling scheme)
  rolled_data <- data[l:(W + l - 1), ]

  # 2. Estimate TR on split data, using whatever formula is desired
  TR_estimate <- lm(formula, data = rolled_data)

  # 3. Pull out coefficients & compute confidence intervals
  all_coefs <- coef(TR_estimate)
  all_cis <- confint(TR_estimate)


  TR_roll[l, var_names] <- all_coefs[var_names]
  TR_roll[l, lower_col_names] <- all_cis[var_names, 1]
  TR_roll[l, upper_col_names] <- all_cis[var_names, 2]  }

# Plotting (note: this part is not modular, obviously)
plot_rolling_coefs(TR_roll, "rate_lag", var_name_title="Rate Lag")

```

Rolling Coefficient Estimate: Rate Lag


with 95% confidence interval



```
plot_rolling_coefs(TR_roll, "inflation_gap", var_name_title="Inflation (Gap)")
```

Rolling Coefficient Estimate: Inflation (Gap)


with 95% confidence interval



```
plot_rolling_coefs(TR_roll, "output_gap", var_name_title="Output Gap")
```

Rolling Coefficient Estimate: Output Gap

with 95% confidence interval



Forecasting Model Evaluation

Helpers

```
#-
# Helper function for adding p-value significance stars
#-

format_p_values_with_stars <- function(p) {
  stars <- case_when(
    p < 0.01 ~ "***",
    p < 0.05 ~ "**",
    p < 0.10 ~ "*",
    TRUE      ~ "")
  paste0(format(round(p, 4), nsmall = 3), " ", stars)}

#-
# HELPER FUNCTION FOR MINCER-ZARNOWITZ REPORTING
#-

# This function runs the Mincer-Zarnowitz regression (Actuals ~ Forecasts)
# for each horizon h and tests the joint null hypothesis H0: (alpha, beta) = (0, 1).
#-

generate_mincer_zarnowitz_report <- function(F_model,
                                             Actual_values,
                                             H,
                                             model_caption,
                                             format = "html") {

  # Pre-allocate storage for results
  mz_results <- data.frame(
    Horizon = 1:H,
    Alpha = numeric(H),
    Beta = numeric(H),
    P_Value_Joint_Test = numeric(H))

  for (h in 1:H) {
    # 1. Create a clean data frame for this horizon
    #   This pairs the forecasts and actual values and removes any NAs,
    #   ensuring they remain perfectly aligned.
    df_h <- data.frame(
      actuals = Actual_values[[h]],
      forecasts = F_model[[h]] ) %>%
      na.omit()

    # Check if we have enough data to run the regression (at least 2 obs)
    if (nrow(df_h) > 2) {
      # 2. Run MZ regression
      mz_reg <- lm(actuals ~ forecasts, data = df_h)

      # 3. Get coefficients
      coeffs <- summary(mz_reg)$coefficients
      mz_results$Alpha[h] <- coeffs[1, 1]
      mz_results$Beta[h]  <- coeffs[2, 1]
```

```

# Using NW errors as seen in class, with lag selection h-1
v_matrix <-
  if (h == 1) {
    # h=1: No autocorrelation, use standard "White" (HC) errors
    sandwich::vcovHC(mz_reg, type = "HC3")
  } else {
    # h>1: Use Newey-West, manually setting lag = h-1
    sandwich::NeweyWest(mz_reg, lag = h - 1)}

# 4. Test Joint Hypothesis H0: Alpha = 0 AND Beta = 1 and store pvalues
test_joint <- linearHypothesis(mz_reg,
                                c("(Intercept) = 0", "forecasts = 1"), vcov. = v_matrix)
mz_results$P_Value_Joint_Test[h] <- test_joint$"Pr(>F)"[2]

} else {
  # Not enough data to run regression for this horizon
  mz_results$Alpha[h] <- NA_real_
  mz_results$Beta[h] <- NA_real_
  mz_results$P_Value_Joint_Test[h] <- NA_real_ }

# Format the results for the table
mz_results <- mz_results %>%
  mutate(Alpha = round(Alpha, 4),
        Beta = round(Beta, 4),
        P_Value_Joint_Test = format_p_values_with_stars(P_Value_Joint_Test))

# Create the table
table_output <- kable(
  mz_results,
  format = format,
  booktabs = TRUE,
  caption = model_caption,
  digits = 4,
  col.names = c("h", "Alpha", "Beta", "pv(Joint)"),
  escape = FALSE ) %>%
  kable_styling(
    latex_options = c("striped", "scale_down"),
    position = "center") %>%
  column_spec(1, bold = TRUE, border_right = TRUE) %>%
  column_spec(4, monospace = TRUE) %>%
  footnote(
    general = "pv(Joint) is the p-value for the joint hypothesis H_0: (Alpha, Beta) = (0, 1). A high p-value indicates that the null hypothesis cannot be rejected at the specified significance level. The significance levels are indicated by the asterisks: ***, **, *, and .",
    symbol = c(
      "Signif. codes:  '***' 0.01,  '**' 0.05,  '*' 0.1"),
    general_title = "Note:",
    symbol_title = "",
    footnote_as_chunk = TRUE,
    threeparttable = TRUE)
  return(table_output) }

#-----

```

```

# HELPER FUNCTION FOR REPORTING DM tests
#-----

# This function creates the DM tests and kable output
generate_report_table <- function(FE_TR_model, FE_BM_model, H, model_caption, format = "html") { MSFE_T
  MSFE_BM = numeric(H)

  # Calculate MSFES
  for (h in 1:H) {
    # Ensure errors are cleaned of NAs
    fe1 <- na.omit(FE_TR_model[[h]])
    fe2 <- na.omit(FE_BM_model[[h]])

    MSFE_TR[h] = mean((fe1)^2)
    MSFE_BM[h] = mean((fe2)^2)}

  # Run DM Tests
  DMpvalues = matrix(, nrow = H, ncol = 3)
  colnames(DMpvalues) <- c("DM_Two_Sided", "DM_Greater", "DM_Lesser")
  for (h in 1:H){
    # Note: dm.test needs the *full* (un-omitted) error vectors
    # to align them properly, hence using the original list inputs
    x1 = dm.test(e1 = FE_BM_model[[h]], e2 = FE_TR_model[[h]], h = h)
    x2 = dm.test(e1 = FE_BM_model[[h]], e2 = FE_TR_model[[h]], h = h, alternative = "greater")
    x3 = dm.test(e1 = FE_BM_model[[h]], e2 = FE_TR_model[[h]], h = h, alternative = "less")
    DMpvalues[h, 1] = round(x1$p.value, digits = 4)
    DMpvalues[h, 2] = round(x2$p.value, digits = 4)
    DMpvalues[h, 3] = round(x3$p.value, digits = 4)}

  # Create final table data
  forecast_comparison <- data.frame(
    Horizon = 1:H,
    MSFE_TR = MSFE_TR,
    MSFE_BM = MSFE_BM) %>%
    mutate(Ratio_TR_vs_BM = MSFE_TR / MSFE_BM)

  forecast_comparison <- bind_cols(forecast_comparison, as.data.frame(DMpvalues))

  final_data_formatted <- forecast_comparison %>%
    mutate(across(starts_with("DM_"), format_p_values_with_stars))

  # Create the kable table
  table_output <- kable(
    final_data_formatted,
    format = format,
    booktabs = TRUE,
    caption = model_caption,
    digits = 4,
    col.names = c("h", "MSFE TR", "MSFE BM", "Ratio", "DM Two-Sided", "DM Greater", "DM Lesser"),
    escape = FALSE) %>%
    kable_styling(
      latex_options = c("striped", "scale_down"),
      position = "center") %>%
}

```

```

column_spec(1, bold = TRUE, border_right = TRUE) %>%
column_spec(5:7, monospace = TRUE) %>%
footnote(
  general = "TR refers to the forecast made with an estimated Taylor Rule. BM refers to a benchmark symbol = c(
    "'DM Greater' tests if the TR model is significantly more accurate than the BM model.",
    "'DM Lesser' tests if the TR model is significantly less accurate than the BM model."),
  general_title = "Note:",
  symbol_title = "DM Test Alternative Hypotheses (H_A):",
  footnote_as_chunk = TRUE,
  threeparttable = TRUE)
return(table_output)}

```

Estimation

Pseudo-out of sample recursive estimation scheme of direct forecasts for all Taylor Rule formulas and a benchmark ARIMA specification.

```

#parameters
R = 85 # Chow: Structural breaks at R=55 and R=85
cat("Evaluation sample starts after ",as.character(data$quarter[R]), ". ",sep="")

```

Evaluation sample starts after 2020 Q1.

```

P = nrow(data) - R #but will effectively be: P = T-h-R
H = 10 #number of different horizons (takes 10 to go until 2027 Q4)

#note: we are doing a recursive estimation scheme for out-of-sample tests
#note: we are doing direct forecasts

#-----
# 1. DEFINE THE TAYLOR RULE (TR) MODEL FORMULAS
#-----

# TR specifications (using either current inflation or inflation expectations
# according to configuration, same with HP vs Hamilton)
formula_1 <- rate ~ inflation_gap + output_gap
formula_2 <- shadowrate ~ inflation_gap + output_gap
formula_3 <- rate ~ rate_lag + inflation_gap + output_gap
formula_4 <- shadowrate ~ shadowrate_lag + inflation_gap + output_gap

#-----
# 2. PRE-ALLOCATE STORAGE FOR ALL RESULTS
#-----

# We need 4 lists for the TR models, 1 list for the shared benchmark
init_storage_list <- function(H, P) {
  storage <- vector("list", length = H)
  for (h in 1:H) {
    storage[[h]] <- rep(NA_real_, P)
  }
  return(storage)}

```

```

# Storage for realised values
Actuals <- init_storage_list(H, P)

# Storage for Forecasts
F_TR_1 <- init_storage_list(H, P) # Model 1: shadowrate, no lag
F_TR_2 <- init_storage_list(H, P) # Model 2: rate, no lag
F_TR_3 <- init_storage_list(H, P) # Model 3: shadowrate, with lag
F_TR_4 <- init_storage_list(H, P) # Model 4: rate, with lag
F_BM   <- init_storage_list(H, P) # Benchmark: ARIMA

# Storage for Forecast Errors
FE_TR_1 <- init_storage_list(H, P) # Model 1: shadowrate, no lag
FE_TR_2 <- init_storage_list(H, P) # Model 2: rate, no lag
FE_TR_3 <- init_storage_list(H, P) # Model 3: shadowrate, with lag
FE_TR_4 <- init_storage_list(H, P) # Model 4: rate, with lag
FE_BM   <- init_storage_list(H, P) # Benchmark: ARIMA

#-----
# 3. SETUP & RUN THE PARALLEL BACKTESTING LOOP
#-----

num_cores <- detectCores() - 1
cl <- makeCluster(num_cores)
registerDoParallel(cl)

# .export sends read-only objects to each core
# .packages loads libraries on each core
worker_results <- foreach(
  p = P:1,
  .packages = c("forecast", "stats", "dplyr"),
  .export = c("data", "H", "formula_1", "formula_2", "formula_3", "formula_4")
) %dopar% {

  # 1. Define splits (with rolling scheme)
  training <- data[(1 + nrow(data) - R - p):(nrow(data) - p), ]
  testing <- data[nrow(data) - (p - 1)):nrow(data), ]

  # --- 2. Fit common models only once ---
  # note: d=1 for interest and inflation as non-stationary
  inflation_arma <- my.auto.arima(training$inflation_gap, max.p=4, max.q=4, d=1)
  outputgap_arma <- my.auto.arima(training$output_gap, max.p=4, max.q=4, d=0)
  interest_arma <- my.auto.arima(training$rate, max.p=4, max.q=4, d=1) # Benchmark

  # --- 3. Get common forecasts only once (all H horizons) ---
  inflation_forecasts <- my.forecast(inflation_arma, h = H)
  outputgap_forecasts <- my.forecast(outputgap_arma, h = H)
  BMpredicted_rates <- my.forecast(interest_arma, h = H)

  # --- 4. Fit the 4 TR models ---
  TR_model_1 <- lm(formula_1, data = training)
  TR_model_2 <- lm(formula_2, data = training)
  TR_model_3 <- lm(formula_3, data = training)
  TR_model_4 <- lm(formula_4, data = training)

```

```

# --- 5. Build forecast input data & get forecasts for non-lagged models ---
# These are direct forecasts
new_data_base <- data.frame(
  inflation_gap = inflation_forecasts,
  output_gap = outputgap_forecasts)

TR_preds_1 <- round(pmax(predict(TR_model_1, new_data_base), min(data$rate)) / 0.25) * 0.25
TR_preds_2 <- round(pmax(predict(TR_model_2, new_data_base), min(data$rate)) / 0.25) * 0.25
BM_preds <- round(pmax(BMpredicted_rates, min(data$rate)) / 0.25) * 0.25

# --- 6. Get forecasts for lagged models via iteration ---
# We must loop 1 step at a time, feeding forecasts back in.

# a) Pre-allocate storage for H forecasts
TR_preds_3 <- numeric(H)
TR_preds_4 <- numeric(H)

# b) Get the last known lag from the training set (lag for h=1 forecast)
current_rate_lag <- last(training$rate)
current_shadowrate_lag <- last(training$shadowrate)

# Loop for iterative forecasting
for (h in 1:H) {
  # --- Prepare dataset for predictions ---
  new_data_3_h <- data.frame(
    inflation_gap = inflation_forecasts[h],
    output_gap = outputgap_forecasts[h],
    rate_lag = current_rate_lag)
  new_data_4_h <- data.frame(
    inflation_gap = inflation_forecasts[h],
    output_gap = outputgap_forecasts[h],
    shadowrate_lag = current_shadowrate_lag)

  # Get the forecast values (keep for lag, and then round for actual prediction)
  pred_3_h <- predict(TR_model_3, new_data_3_h)
  TR_preds_3[h] <- round(pmax(pred_3_h, min(data$rate)) / 0.25) * 0.25
  pred_4_h <- predict(TR_model_4, new_data_4_h)
  TR_preds_4[h] <- round(pmax(pred_4_h, min(data$rate)) / 0.25) * 0.25

  # Update lag for h+1
  current_shadowrate_lag <- pred_4_h
  current_rate_lag <- pred_3_h }

# --- 7. Get actual values in evaluation sample ---
actual_rates <- testing$rate[1:H]

# --- 8. Return all FORECASTS and ACTUALS from the worker ---
list(F_TR_FORMULA_1 = TR_preds_1,
     F_TR_FORMULA_2 = TR_preds_2,
     F_TR_FORMULA_3 = TR_preds_3,
     F_TR_FORMULA_4 = TR_preds_4,
     F_BM = BM_preds,
     actuals = actual_rates) }

```

```

# --- Stop the Cluster ---
stopCluster(cl)
rm(cl)

#-----
# 4. UNPACK PARALLEL RESULTS INTO STORAGE LISTS
#-----

# 'worker_results' is a list of P lists. We need to re-organize it.
for (i in 1:P) {
  # i=1 corresponds to p=P, i=2 to p=P-1, ... i=P to p=1
  # This 'storage_index' matches the loop order
  storage_index <- i
  p_results <- worker_results[[i]]

  for (h in 1:H) {
    # Get the raw values for this h
    actual_val <- p_results$actuals[h]
    f_tr1_val  <- p_results$F_TR_FORMULA_1[h]
    f_tr2_val  <- p_results$F_TR_FORMULA_2[h]
    f_tr3_val  <- p_results$F_TR_FORMULA_3[h]
    f_tr4_val  <- p_results$F_TR_FORMULA_4[h]
    f_bm_val   <- p_results$F_BM[h]

    # Store Actuals (for MZ)
    Actuals[[h]][storage_index] <- actual_val

    # Store Forecasts (for MZ)
    F_TR_1[[h]][storage_index] <- f_tr1_val
    F_TR_2[[h]][storage_index] <- f_tr2_val
    F_TR_3[[h]][storage_index] <- f_tr3_val
    F_TR_4[[h]][storage_index] <- f_tr4_val
    F_BM[[h]][storage_index]   <- f_bm_val

    # Calculate and Store Errors (for MSFE/DM)
    FE_TR_1[[h]][storage_index] <- f_tr1_val - actual_val
    FE_TR_2[[h]][storage_index] <- f_tr2_val - actual_val
    FE_TR_3[[h]][storage_index] <- f_tr3_val - actual_val
    FE_TR_4[[h]][storage_index] <- f_tr4_val - actual_val
    FE_BM[[h]][storage_index]   <- f_bm_val - actual_val } }

#-----
# 5. RENDER RESULTS MORE INTUITIVE FOR FURTHER ANALYSIS
#-----


# Convert the forecast lists (F_TR_x) into single dataframes
forecast_to_df <- function(forecast_list, period) {
  # Convert each element to numeric (benchmark is ts object, which is bad)
  numeric_list <- lapply(forecast_list, function(x) as.numeric(x)) #just make each list inside numeric
  df <- as.data.frame(numeric_list)
  # Add period and horizon
  df$period <- period #first list is all horizon 1 forecasts, gives this to all observations
  df$horizon <- 1:nrow(df) #counts rows and gives each the horizon corresponding to it
}

```

```

df}

# Apply to all forecasting models
eval_all_models <- do.call(rbind, lapply(seq_along(worker_results), function(i) {
  forecast_to_df(worker_results[[i]], period = i)))
eval_all_models[] <- lapply(eval_all_models, function(x) as.numeric(x))

```

Spaghetti Plots

```

# Select the model to plot based on initial option
if (USE_FORMULA == "Formula 1") {
  model <- "F_TR_FORMULA_1"
  model_formula <- formula_1
  model_name <- "Taylor Rule Formula 1"
} else if (USE_FORMULA == "Formula 2") {
  model <- "F_TR_FORMULA_2"
  model_formula <- formula_2
  model_name <- "Taylor Rule Formula 2"
} else if (USE_FORMULA == "Formula 3") {
  model <- "F_TR_FORMULA_3"
  model_formula <- formula_3
  model_name <- "Taylor Rule Formula 3"
} else if (USE_FORMULA == "Formula 4") {
  model <- "F_TR_FORMULA_4"
  model_formula <- formula_4
  model_name <- "Taylor Rule Formula 4" }


# period = date the forecast was made
# date_of_forecast = the future date we are predicting
eval_all_models$date_of_forecast <- eval_all_models$period + eval_all_models$horizon

# Spaghetti plot with color per period
ggplot(eval_all_models, aes(x = date_of_forecast, y = .data[[model]], group = period, color = factor(pe
  geom_line(alpha = 0.7) +      # forecast lines
  geom_point(shape = 2, alpha = 0.7) +
  
  # Actuals as black baseline
  geom_line(aes(y = actuals), color = "black", size = 1) +
  geom_point(aes(y = actuals), color = "black", shape = 4, alpha = 0.5) +
  labs(title = "Evaluation Sample Forecasts vs Realised Values",
       x = "Period",
       y = "Interest Rate",
       color = "Forecast Origin",
       subtitle = paste("For model based on:", model_name)) +
  theme_minimal() +
  theme(plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
        plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))

```

Evaluation Sample Forecasts vs Realised Values

For model based on: Taylor Rule Formula 3




Plots of FE

```
# Compute forecast error
eval_all_models$forecast_error <- eval_all_models[[model]] - eval_all_models$actuals

#compute date_of_forecast for x-axis
eval_all_models$date_of_forecast <- eval_all_models$period + eval_all_models$horizon

ggplot(eval_all_models, aes(x = date_of_forecast, group = period)) +
  # Forecast error lines
  geom_line(aes(y = forecast_error), color = "blue", alpha = 0.5) +
  geom_point(aes(y = forecast_error), color = "blue", alpha = 0.5, shape = 4) +
  
  # Actual rates line
  geom_line(aes(y = actuals), color = "black", size = 1) +
  geom_point(aes(y = actuals), color = "black", shape = 4, alpha = 0.5) +
  labs(title = "Evaluation Sample Forecast Errors",
       x = "Below Zero = forecast was too low ; Above Zero = forecast too high",
       y = "Interest Rate and Forecast Error (deviation from interest rate)",
       subtitle = paste("For model based on:", model_name)) +
  theme_minimal() +
  theme(plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
        plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
```

```
axis.title = element_text(size = 12),
axis.text = element_text(size = 10))
```



Density of FE


```
# Version 1: Non-Adjusted Scales
plot_facet <- ggplot(eval_all_models %>% filter(horizon <= H), aes(x = forecast_error)) +
  geom_density(fill = "#3498db", color = "#2980b9", alpha = 0.6) +
  facet_wrap(~horizon, ncol = 4, labeller = label_both, scales = "free") +
  labs(title = "Density of Forecast Errors by Horizon (Non-Adjusted Scale)",
       subtitle = paste("For model based on:", model_name),
       x = "Forecast Error",
       y = "Density") +
  theme_minimal() +
  theme(plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
        plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10),
        strip.background = element_rect(fill = "#ecf0f1", color = NA),
        strip.text = element_text(face = "bold"))
print(plot_facet)
```

```
## Warning: Removed 45 rows containing non-finite outside the scale range
```

```
## (`stat_density()`).
```

Density of Forecast Errors by Horizon (Non-Adjusted Scale)

For model based on: Taylor Rule Formula 3




```
# Version 2: Adjusted scales
plot_facet <- ggplot(eval_all_models %>% filter(horizon <= H), aes(x = forecast_error)) +
  geom_density(fill = "#3498db", color = "#2980b9", alpha = 0.6) +
  facet_wrap(~horizon, ncol = 4, labeller = label_both) +
  labs(title = "Density of Forecast Errors by Horizon (Adjusted Scale)",
       subtitle = paste("For model based on:", model_name),
       x = "Forecast Error",
       y = "Density") +
  theme_minimal() +
  theme(plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
        plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10),
        strip.background = element_rect(fill = "#ecf0f1", color = NA),
        strip.text = element_text(face = "bold"))
print(plot_facet)

## Warning: Removed 45 rows containing non-finite outside the scale range
## (`stat_density()`).
```

Density of Forecast Errors by Horizon (Adjusted Scale)

For model based on: Taylor Rule Formula 3




```
# Version 3: "Ridges"
plot_ridge <- ggplot(eval_all_models %>% filter(horizon <= H),
                      aes(x = forecast_error, y = as.factor(horizon), fill = stat(x))) +
  geom_density_ridges_gradient(scale = 3, rel_min_height = 0.01) +
  scale_fill_viridis_c(name = "Error", option = "C") +
  labs(title = "Density of Forecast Errors by Horizon",
       subtitle = paste("For model based on:", model_name),
       x = "Forecast Error",
       y = "Forecast Horizon") +
  theme_minimal() +
  theme(plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
        plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10))
print(plot_ridge)

## Warning: `stat(x)` was deprecated in ggplot2 3.4.0.
## i Please use `after_stat(x)` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

## Picking joint bandwidth of 0.664
```

Density of Forecast Errors by Horizon

For model based on: Taylor Rule Formula 3




Variance of FE

```
# This gives us the mean forecast error for the h step ahead forecast
var_by_horizon <- eval_all_models %>%
  group_by(horizon) %>%
  summarize(
    mean_fe = mean(forecast_error, na.rm=T),
    var_fe = sd(forecast_error, na.rm=T)^2, n = n() )

ggplot(var_by_horizon, aes(x = horizon, y = var_fe)) +
  geom_line(color = "#2c3e50", size = 1) +
  geom_point(color = "#e74c3c", size = 3, alpha = 0.8) +
  theme_minimal(base_size = 14) +
  scale_y_continuous(labels = scales::comma) +
  labs(title = "Variance of FE by Horizon",
       subtitle = paste("For model based on:", model_name),
       x = "Forecast Horizon",
       y = "Variance of FE") +
  theme(plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
        plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10),
        panel.grid.minor.x = element_blank() )
```

Variance of FE by Horizon

For model based on: Taylor Rule Formula 3



Absolute Performance: Efficiency & Bias

```
# Call MZ-test helper function 4 times.
# Note: These reports is wrapped in tryCatch as it sometimes fails
# If it does fail, simply decrease R in order to have more
# observations, removing potential multicolinearity.

# MZ Report 1: Actual Rate, No Lag
mz_report_1 <- tryCatch({
  generate_mincer_zarnowitz_report(
    F_model = F_TR_1,
    Actual_values = Actuals,
    H = H,
    model_caption = "Mincer-Zarnowitz Test: Actual Rate, No Lag",
    format = format)}, error = function(e) {
  message("Error generating MZ Report (Actual Rate, No Lag): ", e$message)
  message("Skipping this report and continuing...")
  return(NULL)})

# MZ Report 2: Shadow Rate, No Lag (
mz_report_2 <- tryCatch({
  generate_mincer_zarnowitz_report(
    F_model = F_TR_2,
    Actual_values = Actuals,
```

```

H = H,
model_caption = "Mincer-Zarnowitz Test: Shadow Rate, No Lag",
format = format)}, error = function(e) {
message("Error generating MZ Report (Shadow Rate, No Lag): ", e$message)
message("Skipping this report and continuing...")
return(NULL)})

# MZ Report 3: Actual Rate, with Lag
mz_report_3 <- tryCatch({
generate_mincer_zarnowitz_report(
  F_model = F_TR_3,
  Actual_values = Actuals,
  H = H,
  model_caption = "Mincer-Zarnowitz Test: Actual Rate, with Lag",
  format = format)}, error = function(e) {
message("Error generating MZ Report (Actual Rate, with Lag): ", e$message)
message("Skipping this report and continuing...")
return(NULL)}))

# MZ Report 4: Shadow Rate, with Lag
mz_report_4 <- tryCatch({
generate_mincer_zarnowitz_report(
  F_model = F_TR_4,
  Actual_values = Actuals,
  H = H,
  model_caption = "Mincer-Zarnowitz Test: Shadow Rate, with Lag",
  format = format)}, error = function(e) {
message("Error generating MZ Report (Shadow Rate, with Lag): ", e$message)
message("Skipping this report and continuing...")
return(NULL)}))

# MZ Report 5: Benchmark
mz_report_BM <- tryCatch({
generate_mincer_zarnowitz_report(
  F_model = F_BM,
  Actual_values = Actuals,
  H = H,
  model_caption = "Mincer-Zarnowitz Test: Benchmark ARIMA",
  format = format)}, error = function(e) {
message("Error generating MZ Report (Benchmark ARIMA): ", e$message)
message("Skipping this report and continuing...")
return(NULL)}))

list(
  mz_report_1,
  mz_report_2,
  mz_report_3,
  mz_report_4,
  mz_report_BM)

```

[[1]]

[[2]]

[[3]]

Table 9: Mincer-Zarnowitz Test: Actual Rate, No Lag

| h | Alpha | Beta | pv(Joint) |
|-----------|--------|--------|------------|
| 1 | 1.1720 | 0.3204 | 0.2184 |
| 2 | 1.2241 | 0.3856 | 0.7727 |
| 3 | 0.9667 | 0.8076 | 0.8783 |
| 4 | 0.7954 | 1.2455 | 0.2386 |
| 5 | 0.8566 | 1.4292 | 0.0103 ** |
| 6 | 1.0184 | 1.4755 | 0.0000 *** |
| 7 | 1.2284 | 1.3395 | 0.0000 *** |
| 8 | 1.6036 | 1.0913 | 0.0000 *** |
| 9 | 2.2021 | 0.6248 | 0.0001 *** |
| 10 | 2.9629 | 0.0414 | 0.0030 *** |

Note:

pv(Joint) is the p-value for the joint hypothesis $H_0: (\text{Alpha}, \text{Beta}) = (0, 1)$. A high p-value means we fail to reject the null hypothesis of an unbiased, efficient forecast.

* Signif. codes: '***' 0.01, '**' 0.05,
*, 0.1

Table 10: Mincer-Zarnowitz Test: Shadow Rate, No Lag

| h | Alpha | Beta | pv(Joint) |
|-----------|--------|---------|-----------|
| 1 | 1.8201 | -0.3634 | 0.000 *** |
| 2 | 1.8580 | -0.2293 | 0.000 *** |
| 3 | 1.7647 | -0.0424 | 0.000 *** |
| 4 | 1.8174 | 0.0116 | 0.000 *** |
| 5 | 1.9914 | -0.0130 | 0.000 *** |
| 6 | 2.2061 | -0.0425 | 0.000 *** |
| 7 | 2.4557 | -0.0697 | 0.000 *** |
| 8 | 2.7436 | -0.0941 | 0.000 *** |
| 9 | 2.9320 | -0.0671 | 0.000 *** |
| 10 | 3.1602 | -0.0452 | 0.000 *** |

Note:

pv(Joint) is the p-value for the joint hypothesis $H_0: (\text{Alpha}, \text{Beta}) = (0, 1)$. A high p-value means we fail to reject the null hypothesis of an unbiased, efficient forecast.

* Signif. codes: '***' 0.01, '**' 0.05,
*, 0.1

Table 11: Mincer-Zarnowitz Test: Actual Rate, with Lag

| h | Alpha | Beta | pv(Joint) |
|----|--------|---------|------------|
| 1 | 0.0568 | 1.0019 | 0.7940 |
| 2 | 0.2189 | 0.9525 | 0.7692 |
| 3 | 0.4630 | 0.8861 | 0.7036 |
| 4 | 0.7537 | 0.8270 | 0.6691 |
| 5 | 1.1378 | 0.6906 | 0.5505 |
| 6 | 1.4947 | 0.5933 | 0.4541 |
| 7 | 1.9075 | 0.4131 | 0.1574 |
| 8 | 2.3049 | 0.2375 | 0.0148 ** |
| 9 | 2.7110 | 0.0277 | 0.0010 *** |
| 10 | 3.0873 | -0.1496 | 0.0003 *** |

Note:

pv(Joint) is the p-value for the joint hypothesis H_0 : $(\text{Alpha}, \text{Beta}) = (0, 1)$. A high p-value means we fail to reject the null hypothesis of an unbiased, efficient forecast.

* Signif. codes: '***' 0.01, '**' 0.05, '*', 0.1

[[4]]

[[5]]

Relative Performance (against benchmark)

```
# Call DM-test helper function 4 times.

# Report 1: Actual Rate, No Lag
report_1 <- generate_report_table(
  FE_TR_model = FE_TR_1,
  FE_BM_model = FE_BM,
  H = H,
  model_caption = "MSFE Comparison, Trained on Actual Rate, No Lag",
  format = format)

# Report 2: Shadow Rate, No Lag
report_2 <- generate_report_table(
  FE_TR_model = FE_TR_2,
  FE_BM_model = FE_BM,
  H = H,
  model_caption = "MSFE Comparison, Trained on Shadow Rate, No Lag",
  format = format)

# Report 3: Actual Rate, with Lag
report_3 <- generate_report_table(
  FE_TR_model = FE_TR_3,
  FE_BM_model = FE_BM,
  H = H,
```

Table 12: Mincer-Zarnowitz Test: Shadow Rate, with Lag

| h | Alpha | Beta | pv(Joint) |
|-----------|--------|---------|------------|
| 1 | 0.0533 | 0.9137 | 0.0771 * |
| 2 | 0.1926 | 0.7985 | 0.1967 |
| 3 | 0.4409 | 0.6390 | 0.0084 *** |
| 4 | 0.7725 | 0.4865 | 0.0000 *** |
| 5 | 1.1435 | 0.3537 | 0.0000 *** |
| 6 | 1.4715 | 0.2535 | 0.0000 *** |
| 7 | 1.8170 | 0.1656 | 0.0000 *** |
| 8 | 2.2158 | 0.0884 | 0.0000 *** |
| 9 | 2.6636 | 0.0189 | 0.0000 *** |
| 10 | 3.1629 | -0.0418 | 0.0000 *** |

Note:

pv(Joint) is the p-value for the joint hypothesis H_0 : $(\text{Alpha}, \text{Beta}) = (0, 1)$. A high p-value means we fail to reject the null hypothesis of an unbiased, efficient forecast.

* Signif. codes: '***' 0.01, '**' 0.05, *, 0.1

Table 13: Mincer-Zarnowitz Test: Benchmark ARIMA

| h | Alpha | Beta | pv(Joint) |
|-----------|--------|---------|------------|
| 1 | 0.1403 | 0.9439 | 0.4025 |
| 2 | 0.3980 | 0.8509 | 0.1991 |
| 3 | 0.7262 | 0.7193 | 0.2696 |
| 4 | 1.0868 | 0.5906 | 0.3892 |
| 5 | 1.4665 | 0.4339 | 0.2697 |
| 6 | 1.8291 | 0.3006 | 0.1350 |
| 7 | 2.1804 | 0.1495 | 0.3461 |
| 8 | 2.4846 | 0.0308 | 0.0078 *** |
| 9 | 2.7628 | -0.1282 | 0.0000 *** |
| 10 | 2.9597 | -0.3873 | 0.0000 *** |

Note:

pv(Joint) is the p-value for the joint hypothesis H_0 : $(\text{Alpha}, \text{Beta}) = (0, 1)$. A high p-value means we fail to reject the null hypothesis of an unbiased, efficient forecast.

* Signif. codes: '***' 0.01, '**' 0.05, *, 0.1

Table 14: MSFE Comparison, Trained on Actual Rate, No Lag

| h | MSFE TR | MSFE BM | Ratio | DM Two-Sided | DM Greater | DM Lesser |
|-----------|---------|---------|---------|--------------|------------|------------|
| 1 | 4.0655 | 0.1637 | 24.8364 | 0.0001 *** | 0.9999 | 0.0001 *** |
| 2 | 4.0938 | 0.6781 | 6.0369 | 0.0355 ** | 0.9823 | 0.0177 ** |
| 3 | 3.5461 | 1.6776 | 2.1137 | 0.4047 | 0.7977 | 0.2023 |
| 4 | 3.1944 | 2.8924 | 1.1044 | 0.9159 | 0.5421 | 0.4579 |
| 5 | 3.1434 | 4.4596 | 0.7049 | 0.5990 | 0.2995 | 0.7005 |
| 6 | 3.3438 | 5.9961 | 0.5577 | 0.1229 | 0.0614 * | 0.9386 |
| 7 | 3.6083 | 7.7500 | 0.4656 | 0.0120 ** | 0.0060 *** | 0.9940 |
| 8 | 4.1964 | 9.0357 | 0.4644 | 0.0020 *** | 0.0010 *** | 0.9990 |
| 9 | 5.0865 | 10.7260 | 0.4742 | 0.0002 *** | 0.0001 *** | 0.9999 |
| 10 | 6.1823 | 12.0156 | 0.5145 | 0.0000 *** | 0.0000 *** | 1.0000 |

Note: TR refers to the forecast made with an estimated Taylor Rule. BM refers to a benchmark of the interest rate using an ARIMA model. Ratio < 1 indicates that the TR model has lower MSFE. *DM Test Alternative Hypotheses (H_A):*

* 'DM Greater' tests if the TR model is significantly more accurate than the BM model.

† 'DM Lesser' tests if the TR model is significantly less accurate than the BM model.

```
model_caption = "MSFE Comparison, Trained on Actual Rate, with Lag",
format = format)

# Report 4: Shadow Rate, with Lag
report_4 <- generate_report_table(
  FE_TR_model = FE_TR_4,
  FE_BM_model = FE_BM,
  H = H,
  model_caption = "MSFE Comparison, Trained on Shadow Rate, with Lag",
  format = format)

list(report_1, report_2, report_3, report_4)
```

[[1]]

[[2]]

[[3]]

[[4]]

Table 15: MSFE Comparison, Trained on Shadow Rate, No Lag

| h | MSFE TR | MSFE BM | Ratio | DM Two-Sided | DM Greater | DM Lesser |
|----|---------|---------|---------|--------------|------------|------------|
| 1 | 8.8750 | 0.1637 | 54.2182 | 0.0001 *** | 1.0000 | 0.0000 *** |
| 2 | 9.6375 | 0.6781 | 14.2120 | 0.0024 *** | 0.9988 | 0.0012 *** |
| 3 | 9.6678 | 1.6776 | 5.7627 | 0.0371 ** | 0.9815 | 0.0185 ** |
| 4 | 12.3715 | 2.8924 | 4.2773 | 0.2245 | 0.8878 | 0.1122 |
| 5 | 17.0404 | 4.4596 | 3.8211 | 0.3261 | 0.8369 | 0.1631 |
| 6 | 23.9844 | 5.9961 | 4.0000 | 0.3694 | 0.8153 | 0.1847 |
| 7 | 31.9167 | 7.7500 | 4.1183 | 0.3853 | 0.8073 | 0.1927 |
| 8 | 43.1384 | 9.0357 | 4.7742 | 0.3166 | 0.8417 | 0.1583 |
| 9 | 51.8942 | 10.7260 | 4.8382 | 0.2251 | 0.8875 | 0.1125 |
| 10 | 63.5208 | 12.0156 | 5.2865 | 0.0521 * | 0.9739 | 0.0261 ** |

Note: TR refers to the forecast made with an estimated Taylor Rule. BM refers to a benchmark of the interest rate using an ARIMA model. Ratio < 1 indicates that the TR model has lower MSFE. *DM Test Alternative Hypotheses (H_A):*

* 'DM Greater' tests if the TR model is significantly more accurate than the BM model.

† 'DM Lesser' tests if the TR model is significantly less accurate than the BM model.

Table 16: MSFE Comparison, Trained on Actual Rate, with Lag

| h | MSFE TR | MSFE BM | Ratio | DM Two-Sided | DM Greater | DM Lesser |
|----|---------|---------|--------|--------------|------------|-----------|
| 1 | 0.1399 | 0.1637 | 0.8545 | 0.5402 | 0.2701 | 0.7299 |
| 2 | 0.4625 | 0.6781 | 0.6820 | 0.1832 | 0.0916 * | 0.9084 |
| 3 | 1.0099 | 1.6776 | 0.6020 | 0.1328 | 0.0664 * | 0.9336 |
| 4 | 1.6319 | 2.8924 | 0.5642 | 0.1117 | 0.0559 * | 0.9441 |
| 5 | 2.6250 | 4.4596 | 0.5886 | 0.0449 ** | 0.0224 ** | 0.9776 |
| 6 | 3.6172 | 5.9961 | 0.6033 | 0.0143 ** | 0.0072 *** | 0.9928 |
| 7 | 5.0292 | 7.7500 | 0.6489 | 0.0143 ** | 0.0072 *** | 0.9928 |
| 8 | 6.5804 | 9.0357 | 0.7283 | 0.0000 *** | 0.0000 *** | 1.0000 |
| 9 | 8.5817 | 10.7260 | 0.8001 | 0.0220 ** | 0.0110 ** | 0.9890 |
| 10 | 10.5417 | 12.0156 | 0.8773 | 0.1656 | 0.0828 * | 0.9172 |

Note: TR refers to the forecast made with an estimated Taylor Rule. BM refers to a benchmark of the interest rate using an ARIMA model. Ratio < 1 indicates that the TR model has lower MSFE.

DM Test Alternative Hypotheses (H_A): * 'DM Greater' tests if the TR model is significantly more accurate than the BM model.

† 'DM Lesser' tests if the TR model is significantly less accurate than the BM model.

Table 17: MSFE Comparison, Trained on Shadow Rate, with Lag

| h | MSFE TR | MSFE BM | Ratio | DM Two-Sided | DM Greater | DM Lesser |
|----|---------|---------|--------|--------------|------------|-----------|
| 1 | 0.1637 | 0.1637 | 1.0000 | 1.0000 | 0.5000 | 0.5000 |
| 2 | 0.5781 | 0.6781 | 0.8525 | 0.7065 | 0.3533 | 0.6467 |
| 3 | 1.6678 | 1.6776 | 0.9941 | 0.9464 | 0.4732 | 0.5268 |
| 4 | 3.7257 | 2.8924 | 1.2881 | 0.3371 | 0.8315 | 0.1685 |
| 5 | 7.0993 | 4.4596 | 1.5919 | 0.1935 | 0.9033 | 0.0967 * |
| 6 | 11.6367 | 5.9961 | 1.9407 | 0.0913 * | 0.9544 | 0.0456 ** |
| 7 | 18.5458 | 7.7500 | 2.3930 | 0.1470 | 0.9265 | 0.0735 * |
| 8 | 27.9643 | 9.0357 | 3.0949 | 0.1615 | 0.9193 | 0.0807 * |
| 9 | 40.1971 | 10.7260 | 3.7476 | 0.1525 | 0.9238 | 0.0762 * |
| 10 | 55.4635 | 12.0156 | 4.6160 | 0.1044 | 0.9478 | 0.0522 * |

Note: TR refers to the forecast made with an estimated Taylor Rule. BM refers to a benchmark of the interest rate using an ARIMA model. Ratio < 1 indicates that the TR model has lower MSFE.

DM Test Alternative Hypotheses (H_A): * 'DM Greater' tests if the TR model is significantly more accurate than the BM model.

† 'DM Lesser' tests if the TR model is significantly less accurate than the BM model.

Actual Forecast Model

Helpers

```
#-
# Helper function for displaying our final forecast results (table)
#-

display_forecasts <- function(forecast_list,
                                caption = "Interest Rate Forecasts",
                                format = "html") {

  # Determine the number of horizons and corresponding quarters
  H <- length(forecast_list$TR_Forecast)

  forecast_quarters <- seq(from = last(data$quarter) + 0.25,
                           by = 0.25,
                           length.out = H)

  horizon_quarter_label <- paste0(1:H, ":", as.character(forecast_quarters))

  # Create a data frame for display
  forecast_df <- data.frame(
    Horizon_Quarter = horizon_quarter_label,
    Taylor_Rule_Forecast = round(forecast_list$TR_Forecast, 2),
    Benchmark_ARIMA_Forecast = round(forecast_list$BM_Forecast, 2),
    Inflation_Gap_Forecast = round(forecast_list$Inflation_Forecast, 2),
    Output_Gap_Forecast = round(forecast_list$OutputGap_Forecast, 2))

  # Create the table
  table_output <- kable(
    forecast_df,
    format = format,
    digits = 4,
    col.names = c("Horizon: Quarter", "Taylor Rule Forecast", "Benchmark Forecast",
                 "Inflation Forecast", "Output Gap Forecast"),
    caption = caption,
    booktabs = TRUE) %>%
  kable_styling(
    latex_options = "striped",
    position = "center") %>%
  column_spec(1, bold = TRUE, border_right = TRUE)
  return(table_output) }

#-
# Helper function for plotting our final forecast results
#-

plot_forecasts <- function(forecast_list,
                            title = "Interest Rate and Component Forecasts") {
```

```

# Create the data frame for plotting w/ numerical quarters
H <- length(forecast_list$TR_Forecast)
forecast_quarters_yearqtr <- seq(from = last(data$quarter) + 0.25,
                                   by = 0.25,
                                   length.out = H)

# b) numeric version (for plotting)
forecast_quarters_numeric <- as.numeric(forecast_quarters_yearqtr)

# c) character version (for labels)
forecast_quarters_labels <- as.character(forecast_quarters_yearqtr)

forecast_df <- data.frame(
  Quarter = forecast_quarters_numeric,
  "Taylor Rule" = forecast_list$TR_Forecast,
  "Benchmark ARIMA" = forecast_list$BM_Forecast,
  "Inflation" = forecast_list$Inflation_Forecast,
  "Output Gap" = forecast_list$OutputGap_Forecast,
  check.names = FALSE )

# Long format for ggplot
forecast_long <- forecast_df %>%
  pivot_longer(cols = -Quarter,
               names_to = "Forecast_Type",
               values_to = "Value") %>%
  mutate(Plot_Group = case_when(
    Forecast_Type %in% c("Taylor Rule", "Benchmark ARIMA") ~ "Interest Rate Forecasts",
    Forecast_Type %in% c("Inflation", "Output Gap") ~ "Model Input Forecasts"),
  Plot_Group = factor(Plot_Group, levels = c("Interest Rate Forecasts", "Model Input Forecasts")),
  Forecast_Type = factor(Forecast_Type, levels = c("Taylor Rule", "Benchmark ARIMA", "Inflation", "Output Gap")))

# Actual plot
plot <- ggplot(forecast_long, aes(x = Quarter, y = Value, color = Forecast_Type)) +
  geom_line(linewidth = 1.1) +
  geom_point(size = 2.5) +
  facet_wrap(~ Plot_Group, ncol = 1, scales = "free_y") +
  # Prettiness
  labs(title = title,
       x = "Quarter",
       y = "Value (%)",
       color = "Forecast Series",
       subtitle = paste("For model based on:", model_name)) +
  scale_x_continuous(breaks = forecast_quarters_numeric,
                     labels = forecast_quarters_labels) +
  theme_minimal(base_size = 14) +
  theme(legend.position = "bottom",
        plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
        plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10),
        strip.text = element_text(face = "bold", size = 12),
        axis.text.x = element_text(angle = 45, hjust = 1) +
        scale_color_brewer(palette = "Set1")

```

```
    return(plot) }
```

Forecasting

```
# Formula 4 seems to work best
our_predict <- function(data, formula, H){

  # --- 1. Fit inputs and benchmark models ---
  inflation_arma <- my.auto.arima(data$inflation_gap, max.p=4, max.q=4, d=1)
  outputgap_arma <- my.auto.arima(data$output_gap, max.p=4, max.q=4, d=0)
  interest_arma <- my.auto.arima(data$rate, max.p=4, max.q=4, d=1) # Benchmark

  # --- 2. Get forecasts of inputs (all H horizons) ---
  inflation_forecasts <- my.forecast(inflation_arma, h = H)
  outputgap_forecasts <- my.forecast(outputgap_arma, h = H)
  BMpredicted_rates <- my.forecast(interest_arma, h = H)

  # --- 3. Fit TR model
  TR_model <- lm(formula, data = data)

  # --- 4. Build forecast input data frame (iteratively for lags) ---

  # Allocate storage for full horizon
  TR_preds <- numeric(H)

  # Get last known lags (starting point for lagged models)
  current_shadowrate_lag <- last(data$shadowrate)
  current_rate_lag <- last(data$rate)

  for (h in 1:H) {
    new_data_h <- data.frame(
      inflation_gap = inflation_forecasts[h] ,
      output_gap = outputgap_forecasts[h] ,
      shadowrate_lag = current_shadowrate_lag,
      rate_lag = current_rate_lag)

    # Get forecasted values
    pred_h <- predict(TR_model, new_data_h)
    TR_preds[h] <- round(pmax(pred_h, min(data$rate)) / 0.25) * 0.25

    # Update the lag for h+1
    current_rate_lag <- pred_h }

  # --- 5. Compute forecast for BM ---
  BM_preds <- round(pmax(BMpredicted_rates, min(data$rate)) / 0.25) * 0.25

  return(list(TR_Forecast = TR_preds,
              BM_Forecast = BM_preds,
              Inflation_Forecast = inflation_forecasts + 2, #to add back target
              OutputGap_Forecast = outputgap_forecasts ))}
```

Table 18: For model based on: Taylor Rule Formula 3

| Horizon: Quarter | Taylor Rule Forecast | Benchmark Forecast | Inflation Forecast | Output Gap Forecast |
|------------------|----------------------|--------------------|--------------------|---------------------|
| 1: 2025 Q3 | 2.00 | 1.75 | 1.98 | 0.32 |
| 2: 2025 Q4 | 1.75 | 1.25 | 2.04 | 0.19 |
| 3: 2026 Q1 | 1.75 | 1.25 | 1.85 | 0.04 |
| 4: 2026 Q2 | 1.50 | 1.00 | 1.91 | -0.12 |
| 5: 2026 Q3 | 1.50 | 1.00 | 1.95 | -0.21 |
| 6: 2026 Q4 | 1.50 | 1.00 | 1.98 | -0.26 |
| 7: 2027 Q1 | 1.25 | 1.25 | 2.00 | -0.29 |
| 8: 2027 Q2 | 1.25 | 1.25 | 2.01 | -0.28 |
| 9: 2027 Q3 | 1.25 | 1.25 | 2.02 | -0.24 |
| 10: 2027 Q4 | 1.25 | 1.25 | 2.03 | -0.19 |


```
final_forecasts <- our_predict(data = data, formula = model_formula, H = H)
```

```
display_forecasts(final_forecasts,
                  caption = paste("For model based on:", model_name),
                  format = format)
```

```
plot_forecasts(final_forecasts)
```

Interest Rate and Component Forecasts

For model based on: Taylor Rule Formula 3



Prediction Intervals

```
prediction <- var_by_horizon
final_interval <- final_forecasts
prediction$sd_fe <- sqrt(prediction$var_fe)

final_interval$sd <- prediction$sd_fe
final_interval$upper_1_sd <- final_interval$sd + final_interval$TR_Forecast
final_interval$lower_1_sd <- final_interval$sd*(-1) + final_interval$TR_Forecast

final_interval$upper_2_sd <- final_interval$sd*2 + final_interval$TR_Forecast
final_interval$lower_2_sd <- final_interval$sd*2*(-1) + final_interval$TR_Forecast

final_interval <- as.data.frame(final_interval)

ggplot(final_interval, aes(x = seq_len(nrow(final_interval)), y = TR_Forecast, group = 1)) +
  geom_ribbon(aes(ymin = lower_1_sd, ymax = upper_1_sd), fill = "lightgrey", alpha = 0.6) + # interval
  geom_ribbon(aes(ymin = lower_2_sd, ymax = upper_2_sd), fill = "lightgrey", alpha = 0.2) + # interval
  geom_line(color = "#1f78b4", size = 1) + # line
  geom_point(color = "#e31a1c", size = 3) + # points
  labs(title = "Preliminary Forecast plot",
       x = "h step ahead",
       subtitle = paste("For model based on:", model_name)) +
  theme_minimal() +
  theme(plot.title = element_text(face = "bold", size = 14, margin = margin(b=5)),
        plot.subtitle = element_text(size = 12, color = "grey30", margin = margin(b=10)),
        axis.title = element_text(size = 12),
        axis.text = element_text(size = 10),
        axis.text.x = element_text(angle = 45, hjust = 1))
```

Preliminary Forecast plot

For model based on: Taylor Rule Formula 3

